

DIVYA HEGDE

469-592-9542 · divyahegdefte@gmail.com · linkedin.com/in/hegedivya · github.com/divyahegde-07

EDUCATION

Northeastern University

Master of Science in Data Science

Boston

Aug 2025

R V College of Engineering

Bachelor of Engineering in Telecommunication Engineering

Bengaluru

May 2017

SKILLS

Programming & Data: Python, SQL, Apex, Java, R, ETL, Tableau, Airflow, PySpark, CI/CD

AI/ML & Cloud: AWS (SageMaker, Bedrock, Lambda, S3), GCP (Vertex AI), Docker, MLflow, NLP, RAG, LangChain, LangGraph, LangSmith, Vector Databases, MCP, Fine-Tuning (LoRA, PeFT)

EXPERIENCE

The Burnes Center for Social Change

AI Engineer

Boston

Jan 2025 – Aug 2025

- Collaborated with Massachusetts procurement agency to build an LLM-powered contract review tool with RBAC (IAM) on AWS CDK infrastructure, saving 300+ staff hours (\$18,000) first quarter
- Created an LLM-powered contract review tool using AWS Bedrock, OpenSearch for embeddings, and DynamoDB for user sessions, automating vendor contract analysis flagging 80% of non-compliant terms
- Deployed RAG chatbot using Claude 3.5 on AWS with GitHub Actions and CloudFormation, achieving 85% autonomous resolution rate for more than 700 quarterly procurement inquiries and accelerating buying decisions by 65%
- Orchestrated an evaluation pipeline with Step Functions, leveraging RAGAS metrics and expert-in-the-loop feedback, improving chatbot answer accuracy from 71% to 86%

Salesforce

Data Analyst

Hyderabad

Apr 2020 – Jul 2023

- Established data pipeline consolidating fragmented asset management using MuleSoft, SQL, and Tableau, improving asset traceability by 75%, preventing \$1M+ in device loss, and enabling proactive equipment refresh cycles
- Developed incident forecasting model using ARIMA on 12,000+ Salesforce case records, achieving 70% prediction accuracy for P0/P1 incidents with projected 25% reduction in critical outages through proactive detection
- Engineered employee authentication on AWS Connect by integrating Salesforce and Okta, protecting sensitive data of 125,000+ customers and 56,000+ employees
- Volunteered to streamline the application process for Noise Solution, a UK-based nonprofit using music therapy to support children with behavioral challenges, improving app user-friendliness and contributing to 40% reduction in program dropouts

MaxVal IP Services

Software Engineer

Bengaluru

Aug 2018 – Apr 2020

- Led migration of 1M+ customer records using Salesforce Bulk API, SQL, implementing validation rules to ensure 98% transfer success
- Automated post-deployment code sanity checks using Python, eliminating manual process and reducing deployment errors by 90%

PROJECTS

Churn Predictor *Python, Model Training and Serving, Evaluation*

Aug 2025

- Built a churn prediction training pipeline on AWS SageMaker that achieved 84.6% AUC for customer churn prediction, incorporating XGBoost, resulting in a 5.3% performance improvement over baseline models
- Leveraged AWS Glue to transform raw customer data into 30 predictive features, then performed Bayesian optimization on the model
- Deployed model using AWS Model Registry and SageMaker endpoints with packaging for seamless cross-environment deployment

Assistive Buyers' Engine *Python, React, Document Chunking, Agentic AI, Deployment, Evaluation*

Mar 2025

- Built Cognito authenticated AI chatbot using Claude 3.5 Sonnet, DynamoDB for user session management, and an admin dashboard for content governance, reducing inquiry response time by 70%
- Enhanced RAG with tool-calling, conversation memory, and streaming, leveraging Bedrock KB, Amazon Titan embeddings, dynamic document chunking, and confidence-based filtering for highly accurate, context-aware multi-turn dialogues
- Developed React TypeScript frontend and implemented LLM evaluation pipelines for model performance monitoring, user feedback analytics, and automated knowledge base synchronization

Blue Bikes Demand Forecasting *Python, MLOps, VertexAI, Monitoring*

Jan 2025

- Built end-to-end MLOps pipeline for BlueBikes demand prediction using Apache Airflow, integrating ride data with weather information, achieving 89% prediction accuracy while implementing bias mitigation on underrepresented stations
- Implemented CI/CD architecture with GitHub Actions, DVC, and GCP's VertexAI for model registry, enabling automated retraining workflows triggered by data changes and containerized deployment using Docker to Google Cloud Run with Prometheus monitoring

AWARDS

Won 1st place nationally in the 2025 NASPO Academic Collaboration Award for developing AI tools with Massachusetts' OSD